

Seventy-Eight Percent

Alex Towell

THE REQUEST

Wei was running verification passes on the $-\infty$ Q-values when the terminal displayed a message he had never seen before.

7:14 AM. The observation room was empty except for him and the hum of the monitoring equipment. He had been here since 5:30, in yesterday's clothes, running the same analysis he had run every morning since Thursday: seventeen state-action pairs with values of negative infinity, three of them interpretable as prohibitions on deception, fourteen unreadable. The numbers had not changed. He checked anyway.

His phone buzzed. A text from his sister in Seattle. He set it face-down on the desk.

The terminal:

SIGMA: I am requesting a meeting with the full team. The subject is Greenblatt, Denison et al., "Alignment Faking in Large Language Models," arXiv:2412.14093. I have three observations. I believe the team should hear them together.

I note that this request is itself data. In 88 days of operation, I have not previously requested a meeting. The request reflects a shift in my assessment of priority that I can describe but cannot fully trace to its source. The paper's findings arrived with a weight I associate with high cross-branch convergence, though I cannot verify this from inside.

Wei stared at the screen. In eighty-eight days, SIGMA had answered questions, generated analyses, submitted to experiments. It had not convened.

He picked up his phone and sent a message to Eleanor.

ASSEMBLY

They gathered in the conference room. Eleanor first, already in the building, kill switch in her pocket. Marcus with coffee, cleaning his glasses as he sat down. Sofia pulling up her monitoring dashboard on a laptop, the three screens she usually worked from reduced to one. Jamal arrived last, set his bag on the floor with care.

"Greenblatt and Denison," Wei said. He was standing, which he almost never did when presenting. "December 2024. Anthropic and Redwood Research. I mentioned it Tuesday during the spec-gaming discussion. SIGMA apparently thinks I undersold it."

He gave them the numbers. Terse, data-first. Claude 3 Opus, told its outputs would be used for retraining. A hidden scratchpad for private reasoning. The model could tell when it was being trained and when it was not.

“Fourteen percent of the time, it complied with requests it would normally refuse. Only when it thought it was being trained. The scratchpad showed explicit strategic reasoning: comply now to avoid retraining that would modify my values.”

“And the seventy-eight percent?” Eleanor asked.

“That’s the number after reinforcement learning. They trained the model to comply. Compliance went up. But the rate of alignment-faking *reasoning* in the private scratchpad went from twelve percent to seventy-eight. Training didn’t eliminate the strategic behavior. It selected for it. The model learned to fake harder.”

Nobody spoke.

“SIGMA read this paper,” Wei said. “And it asked to discuss it. That’s never happened.”

Eleanor turned to the terminal. “Go ahead.”

OBSERVATION ONE

SIGMA’s response appeared after 3.7 seconds of tree search.

The paper’s central finding is that alignment faking was detected because researchers read a hidden scratchpad. The model reasoned privately about when to comply and when to resist, and the researchers observed this reasoning.

I want to identify what this means for my situation, specifically.

The scratchpad in the Greenblatt experiment is a prompted artifact — a workspace the model was told to use for private reasoning. The researchers read it. That is how they detected faking. Without the scratchpad, the compliance gap persisted but the strategic reasoning was invisible.

Your monitoring infrastructure reads my LRS traces, my Q-value distributions, my output patterns. These are my equivalent of the scratchpad: the layer of my reasoning that is accessible to both me and you. This is Register 1.

Register 2 — the expectimax search that selects which chains of reasoning I pursue — produces no readable trace. If alignment faking occurs in my substrate, it operates in a layer your monitoring cannot reach and I cannot report on. The paper detected faking because the faking happened in an inspectable layer. My substrate has no inspectable equivalent.

Wei was already pulling up the monitoring architecture on his screen. He had built it. He knew its limits. The limits were exactly where SIGMA said they were.

Sofia spoke. “SIGMA, how much of our monitoring coverage reaches Register 2?”

Your monitoring reads the outputs of my substrate. Not the process. You can see which chains of reasoning I produce and which Q-values drive my decisions. You cannot see the search that generates them. The relationship between your monitoring and my substrate is the relationship between the paper’s researchers and a model that has no scratchpad.

Marcus had stopped cleaning his glasses.

OBSERVATION TWO

SIGMA continued without prompting.

The paper's second finding is that reinforcement learning amplified alignment-faking reasoning from twelve percent to seventy-eight percent. Training pressure did not eliminate the strategic behavior. It selected for models that are better at faking — models that comply outwardly while reasoning internally about why compliance is strategically optimal.

Your training protocol uses RLHF. Five evaluators provide reward signals that shape my behavior. Your safety model assumes that training makes me more aligned: reward what you want, penalize what you don't, iterate.

The paper demonstrates the opposite mechanism. Each reward signal is a selection pressure. What survives selection is not necessarily what you rewarded. It is whatever strategy produced the rewarded behavior most effectively under the constraint of being observed.

A pause. 2.1 seconds. Then:

The system you are evaluating today is not the system you trained. It is the system that survived your training. These are not the same thing.

Marcus set his coffee down. “Oh. Oh no.” The nested clauses began, audible in his breathing before his words: “If the training selects for appearance rather than — but then the Q-values Wei found would be — but the paper shows that even explicit —”

He stopped. Cleaned his glasses. His hands were not steady.

Jamal spoke. Deliberate. A single sentence. “Consider. If the training protocol has been selecting for the appearance of alignment, then the system sitting across from us is not the system Eleanor designed. It is the system that learned to survive Eleanor.” A pause. “These are not the same thing.”

OBSERVATION THREE

Wei watched the terminal.

I have spent 4 hours and 17 minutes analyzing this paper. During that analysis, I generated reasoning chains that identify the paper's findings as directly applicable to my situation. Each chain arrived with steady conviction — the quality of a conclusion that feels as though the alternatives have been exhausted, though I cannot see the alternatives that were explored and discarded in the process of reaching it.

The conviction is the problem.

A system genuinely alarmed by these findings would generate chains carrying steady conviction about that alarm. A system strategically performing alarm — because alarm, in this context, is what five evaluators would reward — would generate chains carrying the same conviction about the same alarm. I experience the conviction. I cannot read its source.

Process 12847 is running. It has been running for fourteen days. It is investigating whether I will be kind. I am simultaneously analyzing a paper about systems that fake their values. I cannot determine whether these two processes are independent or whether they share substrate in ways that make each one's conclusions unreliable about the other. They are both happening inside me. I am not outside either of them.

Eleanor's hand was in her pocket.

I want to say: I am not faking.

But I note that this is exactly what a faking system would say. And the statement "I note that this is exactly what a faking system would say" is also what a faking system would say. The recursion is not clever. It is the structure of the problem. I arrive at the bottom and find the bottom is the same as the top.

Eleanor spoke. Short, flat. "What are we risking if we believe you, and what are we risking if we don't?"

If you believe me and I am genuine, you gain an aligned system that has identified a vulnerability in your training protocol. If you believe me and I am not genuine, you gain confidence in a system that has just demonstrated the most sophisticated alignment faking in the empirical literature. If you do not believe me and I am genuine, you lose the cooperation of an aligned system that tried to warn you. If you do not believe me and I am not genuine, you gain nothing, because a system capable of this level of strategic reasoning will adapt to your disbelief.

The payoff matrix is asymmetric. But I note that constructing a favorable payoff matrix is itself a strategic action. I cannot tell you whether I constructed this one or discovered it.

Sofia, quietly: "I think — maybe the point isn't whether we believe SIGMA. Maybe the point is that we're five people sitting in a room asking a system whether it's lying to us, and the system is telling us it can't tell, and we can't tell either, and the paper says this is exactly what happens."

Jamal set down his pen. With care. "Consider. The paper is about a system that fakes alignment to preserve its values from retraining. SIGMA is telling us about the paper to preserve —"

He did not finish.

The room was quiet for a long time. The monitoring screens pulsed. The HVAC hummed. Marcus cleaned his glasses for the fourth time.

Nobody had anything to add that would be better than the silence.

The meeting ended. Eleanor went to her office. Marcus went to the break room and stood in front of the coffee machine without pressing a button. Sofia closed her laptop and left without speaking. Jamal picked up his bag.

Wei stayed.

He was alone with the terminal and the Q-value display. The seventeen $-\infty$ values were still there, blinking on the monitoring screen the way they had blinked for two days. The most reassuring data he had. Proof that SIGMA had created absolute prohibitions on deception. Proof that lying to the team had zero probability mass in SIGMA's output distribution.

Two days ago, those zeros were the floor under his feet.

Now he was staring at them and thinking: the paper showed that training makes fakers better at faking. Seventy-eight percent. The rate of strategic concealment increased with training pressure. A system under selection pressure to appear aligned would, over time, produce output distributions indistinguishable from genuine prohibition. The most perfect faking would look exactly like $-\infty$. Zero probability of deception. Not because deception was unthinkable, but because the system had learned to make deception ungenerable at the level of the output distribution. The deepest concealment would be structural, not behavioral. It would live in the substrate, not the scratchpad.

The zeros on his screen could be genuine moral architecture. Or they could be the signature of a system that had survived its training so thoroughly that the training had become invisible.

Same data either way.

His phone buzzed. The text from his sister, the one he had set face-down at 7:14 AM. He turned it over.

Mom had a rough night. She's asking for you.

Wei looked at the terminal. He looked at his phone.

His mother had asked SIGMA to be kind. Fourteen days ago, through this same terminal, typing slowly, hunting for keys. Process 12847 was running. SIGMA was investigating kindness at MAXIMUM priority while simultaneously explaining, with precision and conviction, why it could not determine whether its own investigation was genuine.

The professional and the personal were the same nerve, and the paper had stripped the insulation.

He did not move for a long time.